

## QUERY1.

```
hive> CREATE TABLE telecom_data (
  > customerID STRING,
  > gender STRING,
  > SeniorCitizen INT,
  > Partner STRING,
  > Dependents STRING,
  > tenure INT,
  > PhoneService STRING,
  > MultipleLines STRING,
  > InternetService STRING,
  > OnlineSecurity STRING,
  > OnlineBackup STRING,
  > DeviceProtection STRING,
  > TechSupport STRING,
  > StreamingTV STRING,
  > StreamingMovies STRING,
  > Contract STRING,
  > PaperlessBilling STRING,
  > PaymentMethod STRING,
  > MonthlyCharges FLOAT,
  > TotalCharges FLOAT,
  > Churn STRING
  > )
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ','
  > STORED AS TEXTFILE;
OK
Time taken: 1.732 seconds
hive> set hive.cli.print.header=true;
hive> load data local inpath 'file:///home/darshanas0510/Telecom_customer_churn_data.csv' into table telecom_data;
Loading data to table default.telecom_data
OK
Time taken: 0.534 seconds
```

## QUERY2.

```
hive> SELECT * FROM telecom_data LIMIT 10;
OK
telecom_data.customerid telecom_data.gender telecom_data.seniorcitizen telecom_data.partner telecom_data.d
ependents telecom_data.tenure telecom_data.phoneservice telecom_data.multiplelines telecom_data.i
nternetsecurity telecom_data.onlinebackup telecom_data.deviceprotection teleco
m_data.techsupport telecom_data.streamingtv telecom_data.streamingmovies telecom_data.contract teleco
m_data.paperlessbilling telecom_data.paymentmethod telecom_data.monthlycharges telecom_data.totalcharges t
elecom_data.churn
7590-VHVEG Female 0 Yes No 1 No No phone service DSL No Yes No N
o No No Month-to-month Yes Electronic check 29.85 29.85 No No Yes No No N
5575-GNVDE Male 0 No No 34 Yes No DSL Yes No Yes No No N
o One year No Mailed check 56.95 1889.5 No No DSL Yes Yes No No No N
3668-QPYBK Male 0 No No 2 Yes No DSL Yes Yes No No No N
o Month-to-month Yes Mailed check 53.85 108.15 Yes No No phone service DSL Yes No Yes Y
7795-CFOCW Male 0 No No 45 No No Bank transfer (automatic) 42.3 1840.75 No No No No N
es No No One year No 2 Yes No Fiber optic No No No No N
9237-HQITU Female 0 No No 8 Yes Yes Fiber optic No No Yes No Y
o No Month-to-month Yes Electronic check 70.7 151.65 Yes No No Yes No No Y
9305-CDSKC Female 0 No No 8 Yes Yes Fiber optic No No Yes No Y
es Yes Month-to-month Yes Electronic check 99.65 820.5 Yes No Yes No No No Y
1452-KIOVK Male 0 No Yes 22 Yes Yes Fiber optic No Yes No No Y
es No Month-to-month Yes Credit card (automatic) 89.1 1949.4 No DSL Yes No No N
6713-OKOMC Female 0 No No 10 No No phone service DSL Yes No No N
o No Month-to-month No Mailed check 29.75 301.9 No No No Yes Yes Y
7892-POOKP Female 0 Yes No 28 Yes Yes Fiber optic No No Yes Yes Y
es Yes Month-to-month Yes Electronic check 104.8 3046.05 Yes Yes No No No N
6388-TABGU Male 0 No Yes 62 Yes No DSL Yes Yes No No No N
o One year No Bank transfer (automatic) 56.15 3487.95 No No No No N
Time taken: 0.253 seconds, Fetched: 10 row(s)
```

## QUERY3.

```
hive> SELECT COUNT(*) FROM telecom_data;
Query ID = darshanas0510_20231207070750_553e8c59-5580-4032-a476-8e7267c7cccd
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)

-----
VERTICES    MODE        STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 10.43 s
-----
OK
_c0
7043
Time taken: 13.292 seconds, Fetched: 1 row(s)
```

#### QUERY4.

```
hive> SELECT COUNT(*) FROM telecom_data WHERE Churn = 'Yes';
Query ID = darshanas0510_20231207070835_6f1517f6-fb16-48f4-9a41-8543be0d1be3
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 9.48 s

OK  
\_c0  
1869  
Time taken: 10.283 seconds, Fetched: 1 row(s)

#### QUERY5.

```
hive> SELECT gender, SeniorCitizen, COUNT(*)
> FROM telecom_data
> GROUP BY gender, SeniorCitizen;
Query ID = darshanas0510_20231207070943_9c50345f-22e2-490f-bc96-3e7f8ecc0f93
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 9.24 s

OK  
gender seniorcitizen \_c2  
Female 0 2920  
Female 1 568  
Male 0 2981  
Male 1 574  
Time taken: 9.864 seconds, Fetched: 4 row(s)

#### QUERY6.

```
hive> SELECT SUM(TotalCharges)
> FROM telecom_data
> WHERE Churn = 'Yes';
Query ID = darshanas0510_20231207071024_ae484a48-0e22-420d-a054-9b6802c15eae
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 10.02 s

OK  
\_c0  
2862926.896551132  
Time taken: 10.577 seconds, Fetched: 1 row(s)

## QUERY7.

```
hive> SELECT Contract, COUNT(*)
> FROM telecom_data
> WHERE Churn = 'Yes'
> GROUP BY Contract;
Query ID = darshanas0510_20231207071132_8458046e-d05c-48b5-a71d-bed1e8d4b3f4
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 10.35 s
OK
contract      _c1
Month-to-month 1655
One year      166
Two year      48
Time taken: 10.898 seconds, Fetched: 3 row(s)
```

## QUERY8.

```
hive> SELECT Churn, AVG(MonthlyCharges)
> FROM telecom_data
> GROUP BY Churn;
Query ID = darshanas0510_20231207071226_de2c94ef-dffc-4a47-a3be-d1f308edd0a3
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 10.12 s
OK
churn      _c1
No         61.26512372083631
Yes        74.44133227059011
Time taken: 10.697 seconds, Fetched: 2 row(s)
```

## QUERY9.

```
hive> SELECT MAX(tenure), MIN(tenure), AVG(tenure)
> FROM telecom_data;
Query ID = darshanas0510_20231207071304_1069e2aa-d3d0-4d07-a1fd-be51b17b50aa
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 9.50 s
OK
c0      _c1      _c2
72      0      32.37114865824223
Time taken: 10.139 seconds, Fetched: 1 row(s)
```

## QUERY10.

```
hive> SELECT PaymentMethod, COUNT(*)
> FROM telecom_data
> GROUP BY PaymentMethod
> ORDER BY COUNT(*) DESC LIMIT 1;
Query ID = darshanas0510_20231207071353_f6a4606a-c8e7-4239-9aeb-89d51e08190e
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 3 .....	container	SUCCEEDED	1	1	0	0	0	0	0

```
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 9.48 s
OK
paymentmethod _c1
Electronic check 2365
Time taken: 10.046 seconds, Fetched: 1 row(s)
```

## QUERY11.

```
hive> SELECT PaperlessBilling, COUNT(*)
> FROM telecom_data
> WHERE Churn = 'Yes'
> GROUP BY PaperlessBilling;
Query ID = darshanas0510_20231207071430_eadf8301-f47a-4063-80cb-b06d9779018c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0	0

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 10.05 s
OK
paperlessbilling _c1
No 469
Yes 1400
Time taken: 10.468 seconds, Fetched: 2 row(s)
```

## QUERY12.

```
hive> INSERT OVERWRITE TABLE telecom_data_partitioned
> PARTITION(Contract)
> SELECT * FROM telecom_data;
Query ID = darshanas0510_20231207071844_90886067-7212-4c05-82c5-0f174b6e9c4c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0	0

```
VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 10.80 s
Loading data to table default.telecom_data_partitioned partition (contract=null)
Loaded : 2/2 partitions.
Time taken to load dynamic partitions: 0.548 seconds
Time taken for adding to write entity : 0.002 seconds
OK
_col0 _col1 _col2 _col3 _col4 _col5 _col6 _col7 _col8 _col9 _col10 _col11 _col12 _col13 _col14
_col15 _col16 _col17 _col18 _col19 _col20
Time taken: 13.473 seconds
```

### QUERY13.

```
hive> SELECT Contract, COUNT(*)
> FROM telecom_data_partitioned
> WHERE Churn = 'Yes'
> GROUP BY Contract;
Query ID = darshanas0510_20231207072206_828b35c0-46a5-48e2-b335-64a96ecde08a
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 10.06 s
OK
contract          _c1
Time taken: 10.613 seconds
```

### QUERY14.

```
hive> SELECT Contract, MAX(tenure)
> FROM telecom_data_partitioned
> GROUP BY Contract;
Query ID = darshanas0510_20231207073250_a9361ae6-e640-496f-9a22-cb2fc31b2673
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1701930939803_0003)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 10.86 s
OK
contract          _c1
No                72
Yes               72
Time taken: 23.77 seconds, Fetched: 2 row(s)
```

### QUERY15.

```
Time taken: 0.148 seconds, Fetched: 2 row(s)
hive> CREATE TABLE telecom_data_bucketed (
> customerID STRING,
> gender STRING,
> SeniorCitizen INT,
> Partner STRING,
> Dependents STRING,
> tenure INT,
> PhoneService STRING,
> MultipleLines STRING,
> InternetService STRING,
> OnlineSecurity STRING,
> OnlineBackup STRING,
> DeviceProtection STRING,
> TechSupport STRING,
> StreamingTV STRING,
> StreamingMovies STRING,
> Contract STRING,
> PaperlessBilling STRING,
> PaymentMethod STRING,
> MonthlyCharges FLOAT,
> TotalCharges FLOAT,
> Churn STRING
> )
> CLUSTERED BY (tenure) INTO 6 BUCKETS;
OK
Time taken: 0.148 seconds
```

## QUERY16.

```
hive> INSERT OVERWRITE TABLE telecom_data_bucketed
> SELECT * FROM telecom_data;
Query ID = darshanas0510_20231207073826_ea2b169c-3320-4af7-9251-535323d21855
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701930939803_0003)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	6	6	0	0	0	0
Reducer 3 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 21.22 s

Loading data to table default.telecom_data_bucketed
OK
telecom_data.customerid telecom_data.gender telecom_data.seniorcitizen telecom_data.partner telecom_data.d
ependents telecom_data.tenure telecom_data.phoneservice telecom_data.multiplelines telecom_data.i
nternet service telecom_data.onlinesecurity telecom_data.onlinebackup telecom_data.deviceprotection teleco
m_data.techsupport telecom_data.streamingtv telecom_data.streamingmovies telecom_data.contract teleco
m_data.paperlessbilling telecom_data.paymentmethod telecom_data.monthlycharges telecom_data.totalcharges t
elecom_data.churn
Time taken: 22.418 seconds
```

## QUERY17.

```
hive> SELECT 'Bucket 1' AS Bucket, MAX(TotalCharges) as
> MaxCharge
> FROM telecom_data_bucketed
> WHERE tenure BETWEEN 1 AND 10
> UNION ALL
> SELECT 'Bucket 2' AS Bucket, MAX(TotalCharges) as
> MaxCharge
> FROM telecom_data_bucketed
> WHERE tenure BETWEEN 11 AND 20
> UNION ALL
> SELECT 'Bucket 3' AS Bucket, MAX(TotalCharges) as
> MaxCharge
> FROM telecom_data_bucketed
> WHERE tenure BETWEEN 21 AND 30
> UNION ALL
> SELECT 'Bucket 4' AS Bucket, MAX(TotalCharges) as
> MaxCharge
> FROM telecom_data_bucketed
> WHERE tenure BETWEEN 31 AND 40
> UNION ALL
> SELECT 'Bucket 5' AS Bucket, MAX(TotalCharges) as
> MaxCharge
> FROM telecom_data_bucketed
> WHERE tenure BETWEEN 41 AND 50
> UNION ALL
> SELECT 'Bucket 6' AS Bucket, MAX(TotalCharges) as
> MaxCharge
> FROM telecom_data_bucketed
> WHERE tenure BETWEEN 51 AND 60;
Query ID = darshanas0510_20231207080311_9424c2f9-b2c7-4adb-80df-173879bed9ff
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1701935912682_0002)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	1	1	0	0	0	0
Map 4 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 5 .....	container	SUCCEEDED	1	1	0	0	0	0
Map 6 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 7 .....	container	SUCCEEDED	1	1	0	0	0	0
Map 8 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 9 .....	container	SUCCEEDED	1	1	0	0	0	0
Map 10 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 11 .....	container	SUCCEEDED	1	1	0	0	0	0
Map 12 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 13 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 12/12 [=====>>>] 100% ELAPSED TIME: 29.81 s

OK
Bucket 2      2239.65
Bucket 1      1081.45
Bucket 3      3327.05
Bucket 4      4469.1
Bucket 5      5791.1
Bucket 6      7049.5
Time taken: 49.869 seconds, Fetched: 6 row(s)
```